

# Distributed Systems

José Orlando Pereira

HASLab / Departamento de Informática  
Universidade do Minho



2011/2012

# Case study

- Simple echo server:
  - Echos all bytes received
  - Handle multiple concurrent clients
- Implementation:
  - Listens on port 12345
  - Read and write back

# Sockets in java.net

```
ServerSocket ss=new ServerSocket(12345);

while(true) {
    Socket s=ss.accept();

    InputStream is=s.getInputStream();
    OutputStream os=s.getOutputStream();

    // i/o

    s.close();
}
```

# Sockets in java.nio

```
ServerSocketChannel ss=SelectorProvider.provider().openServerSocketChannel();
ss.socket().bind(new InetSocketAddress(12345));


while(true) {
    SocketChannel s=ss.accept();

    // i/o

    s.close();
}
```


# Buffers in java.nio

- Buffer = Array + Indexes:




# Buffers in java.nio

- Put/read: advances position, sets content

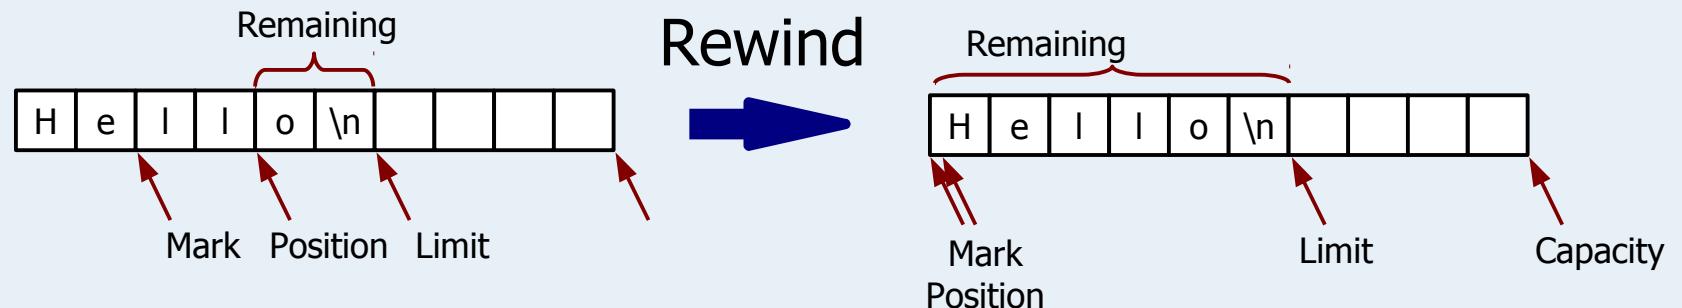
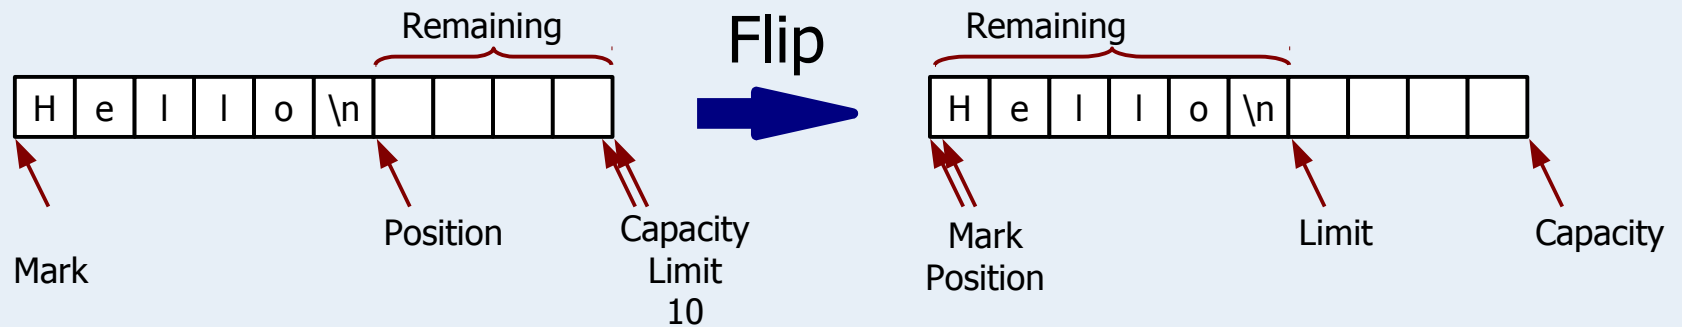



# Buffers in java.nio

- Get/write: advances position, gets content

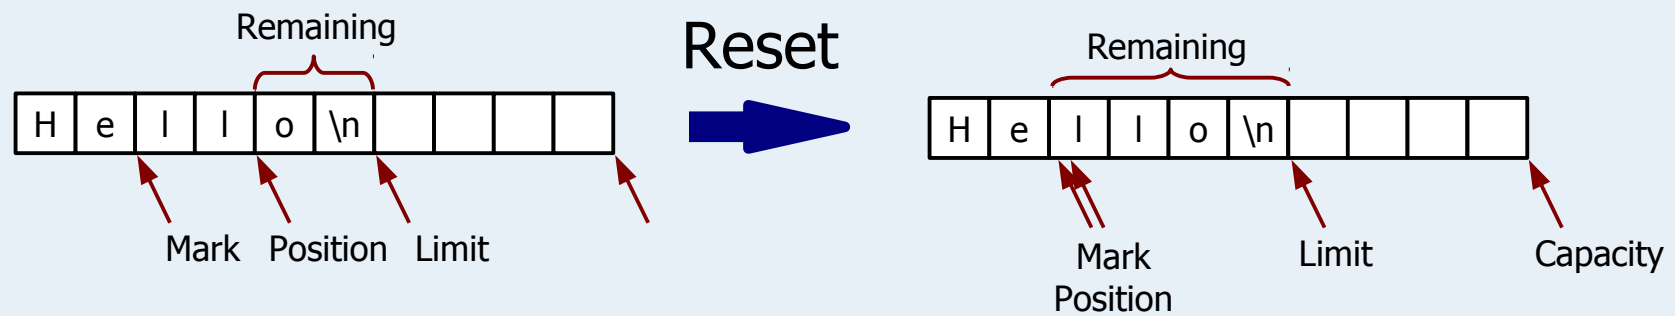



# Buffers in java.nio





# Buffers in java.nio



# Sockets in java.nio

```
ServerSocketChannel ss=SelectorProvider.provider().openServerSocketChannel();
ss.socket().bind(new InetSocketAddress(12345));

while(true) {
    SocketChannel s=ss.accept();

    ByteBuffer buf=ByteBuffer.allocate(100);

    while(s.read(buf)>0) {
        buf.flip();
        s.write(buf);
        buf.clear();
    }

    s.close();
}
```

# Non-blocking I/O

- How to service multiple sockets?
  - Multiple threads
  - Polling with a single thread
- Efficient polling:
  - Use `select()` to wait for I/O
  - Execute I/O operation without blocking

# Non-blocking I/O

```
ServerSocketChannel ss=SelectorProvider.provider().openServerSocketChannel();
ss.socket().bind(new InetSocketAddress(12345));

ss.configureBlocking(false);

Selector sel=SelectorProvider.provider().openSelector();
ss.register(sel, SelectionKey.OP_ACCEPT);

while(true) {
    sel.select();

    for(Iterator<SelectionKey> i=sel.selectedKeys().iterator(); i.hasNext(); ) {
        SelectionKey key = i.next();

        // i/o

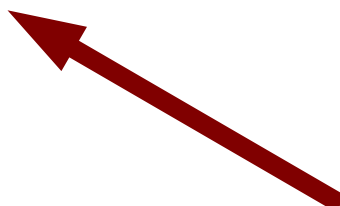
        i.remove();
    }
}
```

# Non-blocking I/O

```
if (key.isAcceptable()) {  
    SocketChannel s=ss.accept();  
  
    s.configureBlocking(false);  
    s.register(sel, SelectionKey.OP_READ);  
}
```

# Non-blocking I/O

```
if (key.isReadable()) {  
    ByteBuffer buf=ByteBuffer.allocate(100);  
    SocketChannel s=(SocketChannel)key.channel();  
  
    int r=s.read(buf);  
    if (r<0) {  
        key.cancel();  
        s.close();  
    } else {  
        buf.flip();  
        s.write(buf);  
    }  
}
```



What if write would block?

# Non-blocking I/O

- Need to poll before writing
- Bytes read must be saved until writing is possible
- Signal interest on writing

```
if (key.isReadable()) {
```

```
    ...
```

```
    } else {
```

```
        buf.flip();
```

```
        key.attach(buf);
```

```
        key.interestOps(SelectionKey.OP_WRITE);
```

```
    }
```

```
}
```



Context

# Non-blocking I/O

- Get bytes attached to key
- Reset interest to reading

```
if (key.isWritable()) {  
    SocketChannel s=(SocketChannel)key.channel();  
    ByteBuffer buf=(ByteBuffer)key.attachment();  
  
    s.write(buf);  
    key.interestOps(SelectionKey.OP_READ);  
}
```



# Object oriented + Event driven

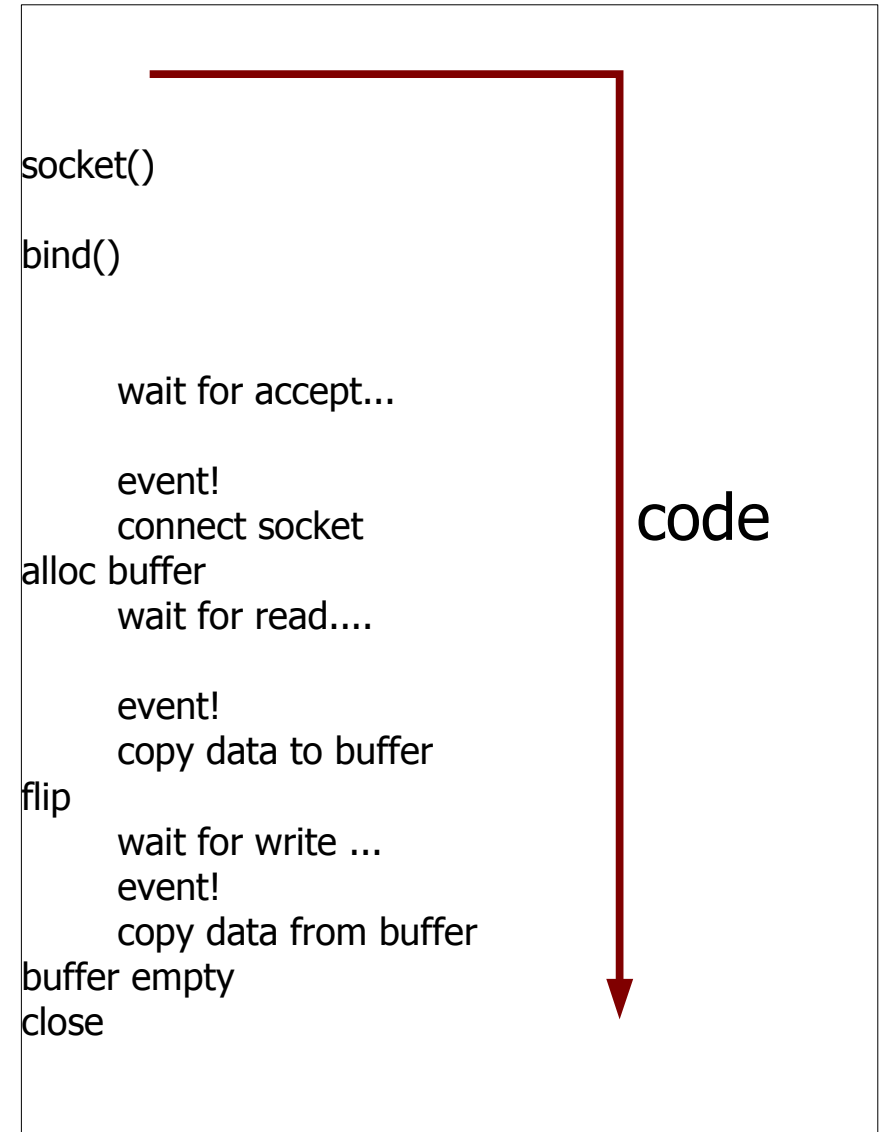
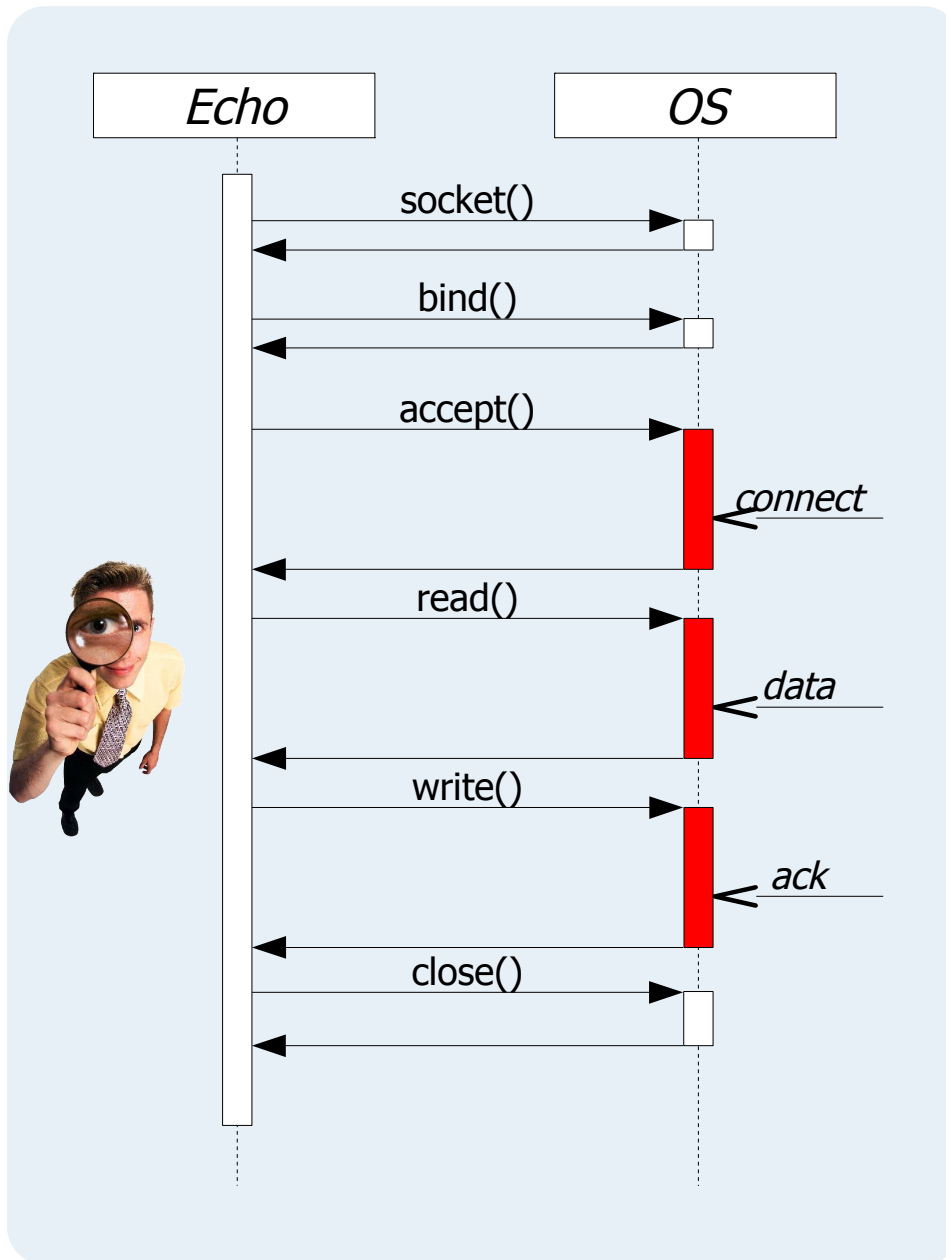
- Encapsulate context data + event-handling code

```
public class Echo {  
    private ByteBuffer buf;  
    // ...  
  
    public Echo(SelectionKey key) {  
        // initialization  
    }  
  
    public void handleRead() throws IOException {  
        // input  
    }  
  
    public void handleWrite() throws IOException {  
        // output  
    }  
}
```

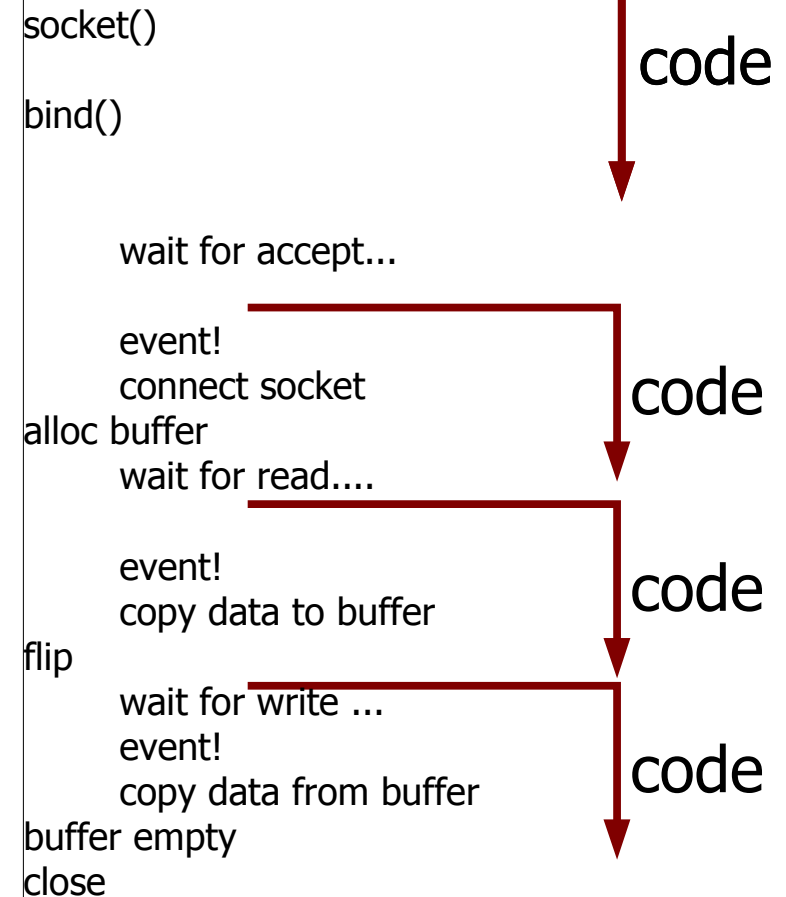
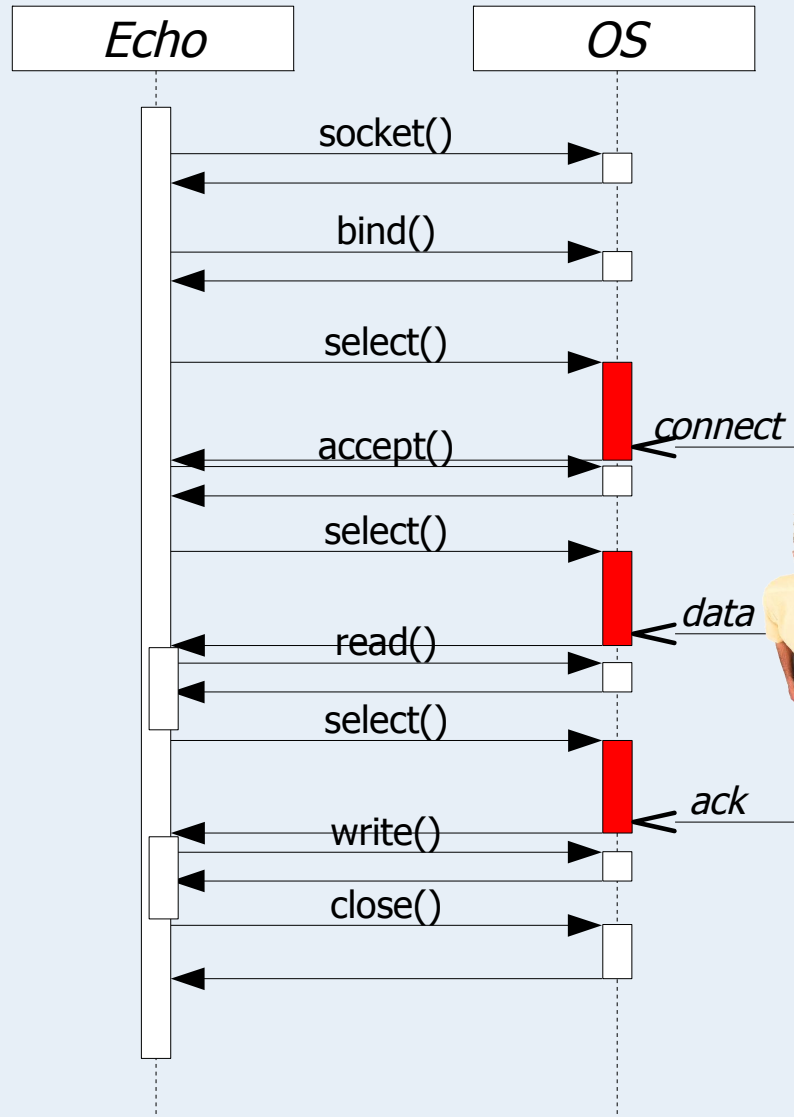
# Object oriented + Event driven

```
if (key.isAcceptable()) {  
    SocketChannel s=ss.accept();  
  
    if (s!=null) {  
        s.configureBlocking(false);  
        SelectionKey nkey=s.register(sel, SelectionKey.OP_READ);  
        nkey.attach(new Echo(nkey));  
    }  
} else if (key.isReadable()) {  
    Echo echo=(Echo)key.attachment();  
    echo.handleRead();  
} else if (key.isWritable()) {  
    Echo echo=(Echo)key.attachment();  
    echo.handleWrite();  
}
```

# Threaded version



# Event-driven version



# Conclusions

- The program can be regarded as either:
  - Being suspended waiting for something to happen
  - Executing some code deterministically without external intervention

# Conclusions

- Both versions of the program:
  - Wait for the same conditions to start executing
  - Execute the same code as a consequence
- If we label such conditions and sections of code equally and log their execution:
  - We cannot distinguish their log files
- If we log the values of key variables in both programs when suspended:
  - We cannot distinguish their log files